

Concurrency Control

CENG315 INFORMATION MANAGEMENT

Three Concurrency Problems

- In a multi-processing environment transactions can interfere with each other.
- Three concurrency problems can arise, that any DBMS must account for and avoid:
 - Lost Update
 - Uncommitted Dependency
 - Inconsistent Analysis




Fig. 16.1 Transaction A loses an update at time $t4$.

The Lost Update Problem

A lost update occurs when a second transaction reads the state of the database prior to the first one writing a change, and then stops on the first one's change with its own update.

The Uncommitted Dependency Problem





Fig. 16.2 Transaction A becomes dependent on an uncommitted change at time t2

An uncommitted dependency occurs when a second transaction relies on a change which has not yet been committed, which is rolled back after the second transaction has begun.



The Inconsistent Analysis Problem

An inconsistent analysis occurs when totals are calculated during interleaved updates.

Fig. 16.4 Transaction A performs an inconsistent analysis

Conflicts

- If A and B are concurrent transactions, problems can occur if A and B want to read or write the same database object, say tuple t .
- There are four possibilities:
 - RR (Read Read): Reads cannot interfere with each other, so there is no problem in this case.
 - RW (Read Write)
 - WR (Write Read)
 - WW (Write Write)

RW Conflict

- *A* reads *t* and then *B* wants to write *t*.
- If *B* is allowed to perform its write, then the inconsistent analysis problem can arise.
 - As we saw in Fig. 16.4.
- Thus, we can say that inconsistent analysis is caused by a RW conflict.

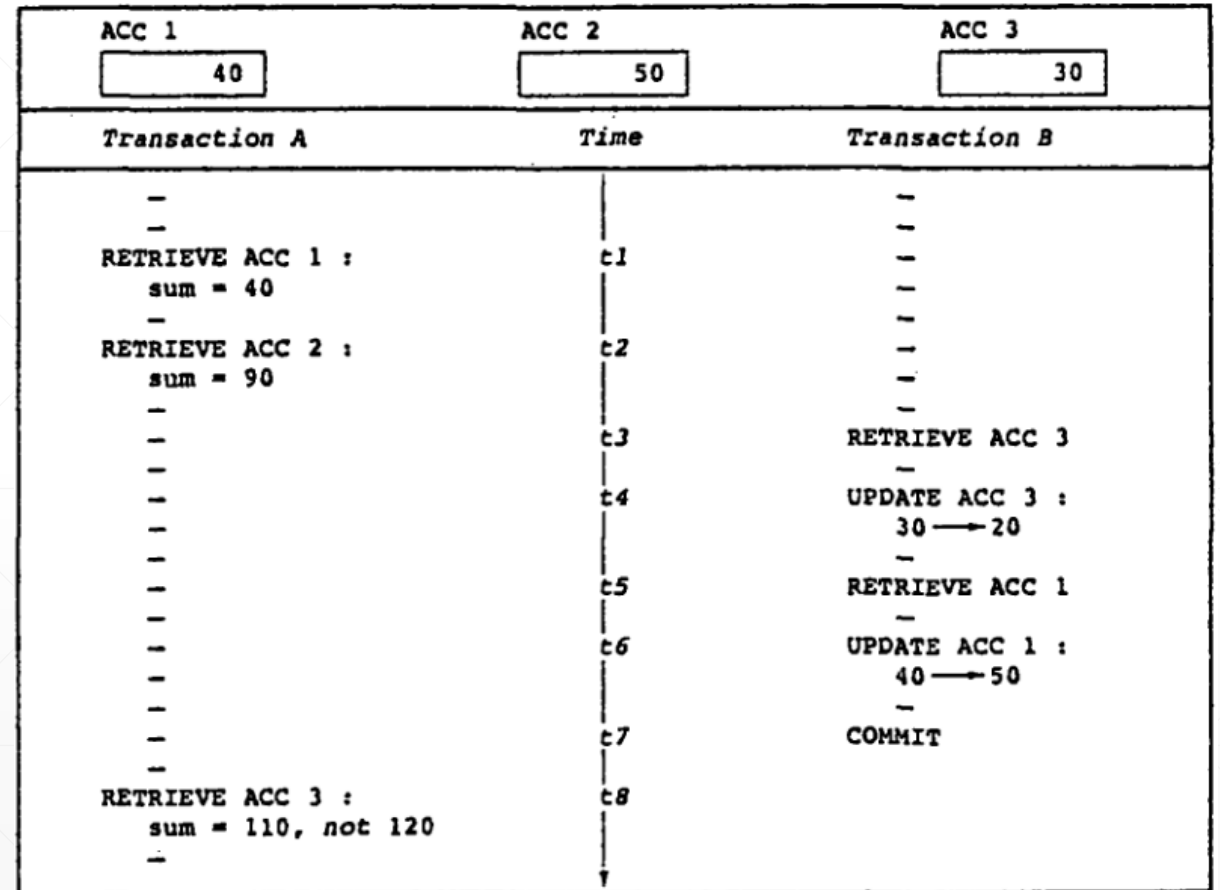


Fig. 16.4 Transaction A performs an inconsistent analysis

WR Conflict

- *B* writes *t* and then *A* wants to read *t*.
- If *A* is allowed to perform its read, then the uncommitted dependency problem can arise.
 - As we saw in Fig. 16.2.




Fig. 16.2 Transaction *A* becomes dependent on an uncommitted change at time *t*2

WR Conflict

- B writes t and then A wants to read t .
- If A is allowed to perform its read, then the uncommitted dependency problem can arise.
 - As we saw in Fig. 16.2.
- Thus, we can say that uncommitted dependencies are caused by WR conflicts.
- Note: A 's read, if it is allowed, is said to be a dirty read.

WW Conflict

- *A* writes *t* and then *B* wants to write *t*.
- If *B* is allowed to perform its write, then the lost update problem can arise.
 - As we saw in Fig. 16.1.




Fig. 16.1 Transaction *A* loses an update at time *t*₄.

WW Conflict

- A writes t and then B wants to write t .
- If B is allowed to perform its write, then the lost update problem can arise.
 - As we saw in Fig. 16.1.
- Thus, we can say that lost updates are caused by WW conflicts.

Locking

- The mentioned problems can all be solved by means of a concurrency control mechanism called locking.
- A transaction locks a portion of the database to prevent concurrency problems.
- Exclusive lock (X) – write lock, will lock out all other transactions
- Shared lock (S) – read lock, will lock out writes, but allow other reads

Lock-Based Protocols

- Lock Compatibility Matrix:

	S	X
S	true	false
X	false	false

- A transaction may be granted a lock on an item if the requested lock is compatible with locks already held on the item by other transactions.
- Any number of transactions can hold shared locks on an item.
- But if any transaction holds an exclusive on the item no other transaction may hold any lock on the item.

Locking Protocol

- A locking protocol is a set of rules that state when a transaction may lock and unlock each of the data items in the database.
- **The Two-Phase Locking Protocol**

Two-Phase Locking Protocol


- A protocol which ensures conflict-serializable schedules.

- Phase 1: **Growing Phase**

- Transaction may obtain locks
 - Transaction may not release locks

- Phase 2: **Shrinking Phase**

- Transaction may release locks
 - Transaction may not obtain locks



Example: Two-Phase Locking Protocol

- Consider the transactions T_{34} and T_{35} :
 - Add lock and unlock instructions to transactions T_{34} and T_{35} , so that they observe the two-phase locking protocol.
 - A transaction requests a shared lock on data item Q by executing the lock-S(Q) instruction.
 - A transaction requests an exclusive lock on data item Q by executing the lock-X(Q) instruction.
 - A transaction can unlock a data item Q by the unlock(Q) instruction.

```
 $T_{34}$ : read(A);  
      read(B);  
      if  $A = 0$  then  $B := B + 1$ ;  
      write(B).
```

```
 $T_{35}$ : read(B);  
      read(A);  
      if  $B = 0$  then  $A := A + 1$ ;  
      write(A).
```


Example: Two-Phase Protocol (Cont.)

- Lock and unlock instructions:

```
T34: read(A);  
      read(B);  
      if A = 0 then B := B + 1;  
      write(B).
```

```
T35: read(B);  
      read(A);  
      if B = 0 then A := A + 1;  
      write(A).
```

```
T34: lock-S(A)  
      read(A)  
      lock-X(B)  
      read(B)  
      if A = 0  
      then B := B + 1  
      write(B)  
      unlock(A)  
      unlock(B)
```

```
T35: lock-S(B)  
      read(B)  
      lock-X(A)  
      read(A)  
      if B = 0  
      then A := A + 1  
      write(A)  
      unlock(B)  
      unlock(A)
```

Deadlock

- Consider the partial schedule:
- Neither T_3 nor T_4 can make progress — executing **lock-S**(B) causes T_4 to wait for T_3 to release its lock on B , while executing **lock-X**(A) causes T_3 to wait for T_4 to release its lock on A .
- Such a situation is called a **deadlock**.
 - To handle a deadlock one of T_3 or T_4 must be rolled back and its locks released.
- The potential for deadlock exists in most locking protocols.

T_3	T_4
lock-X(B) read(B) $B := B - 50$ write(B)	lock-S(A) read(A) lock-S(B)
lock-X(A)	

Starvation

- **Starvation** is also possible if concurrency control manager is badly designed. For example:
 - A transaction may be waiting for an X-lock on an item, while a sequence of other transactions request and are granted an S-lock on the same item.
 - The same transaction is repeatedly rolled back due to deadlocks.
- Concurrency control manager can be designed to prevent starvation.

Example: Two-Phase Protocol and Deadlock

- Can the execution of these transactions result in a deadlock?

```
T34:  lock-S(A)
      read(A)
      lock-X(B)
      read(B)
      if A = 0
      then B := B + 1
      write(B)
      unlock(A)
      unlock(B)
```

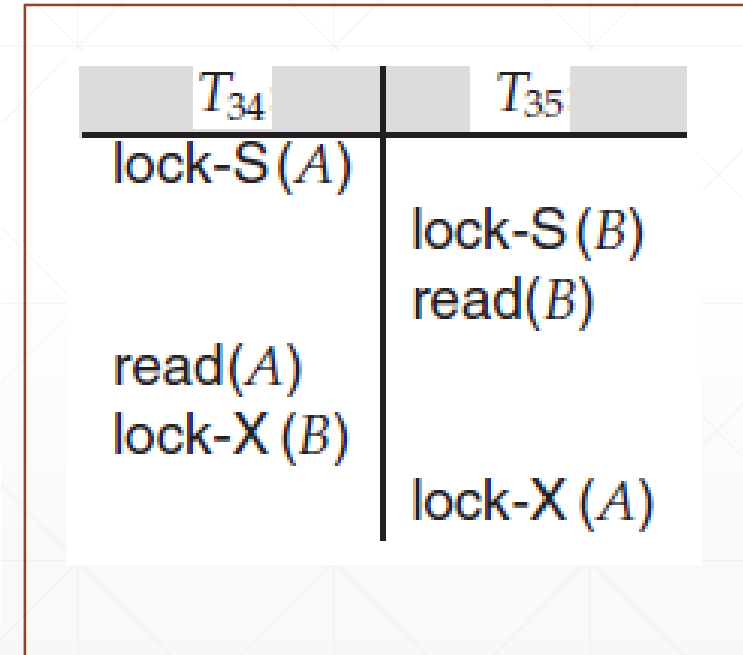
```
T35:  lock-S(B)
      read(B)
      lock-X(A)
      read(A)
      if B = 0
      then A := A + 1
      write(A)
      unlock(B)
      unlock(A)
```

Example: Two-Phase Protocol and Deadlock (Cont.)

- Execution of these transactions can result in deadlock. For example, consider the following partial schedule:

T_{34} :
lock-S(A)
read(A)
lock-X(B)
read(B)
if $A = 0$
then $B := B + 1$
write(B)
unlock(A)
unlock(B)

T_{35} :
lock-S(B)
read(B)
lock-X(A)
read(A)
if $B = 0$
then $A := A + 1$
write(A)
unlock(B)
unlock(A)






Fig. 16.1 Transaction A loses an update at time $t4$.

The Lost Update Problem

A lost update occurs when a second transaction reads the state of the database prior to the first one writing a change, and then stomps on the first one's change with its own update.

The Lost Update Problem - Revisited

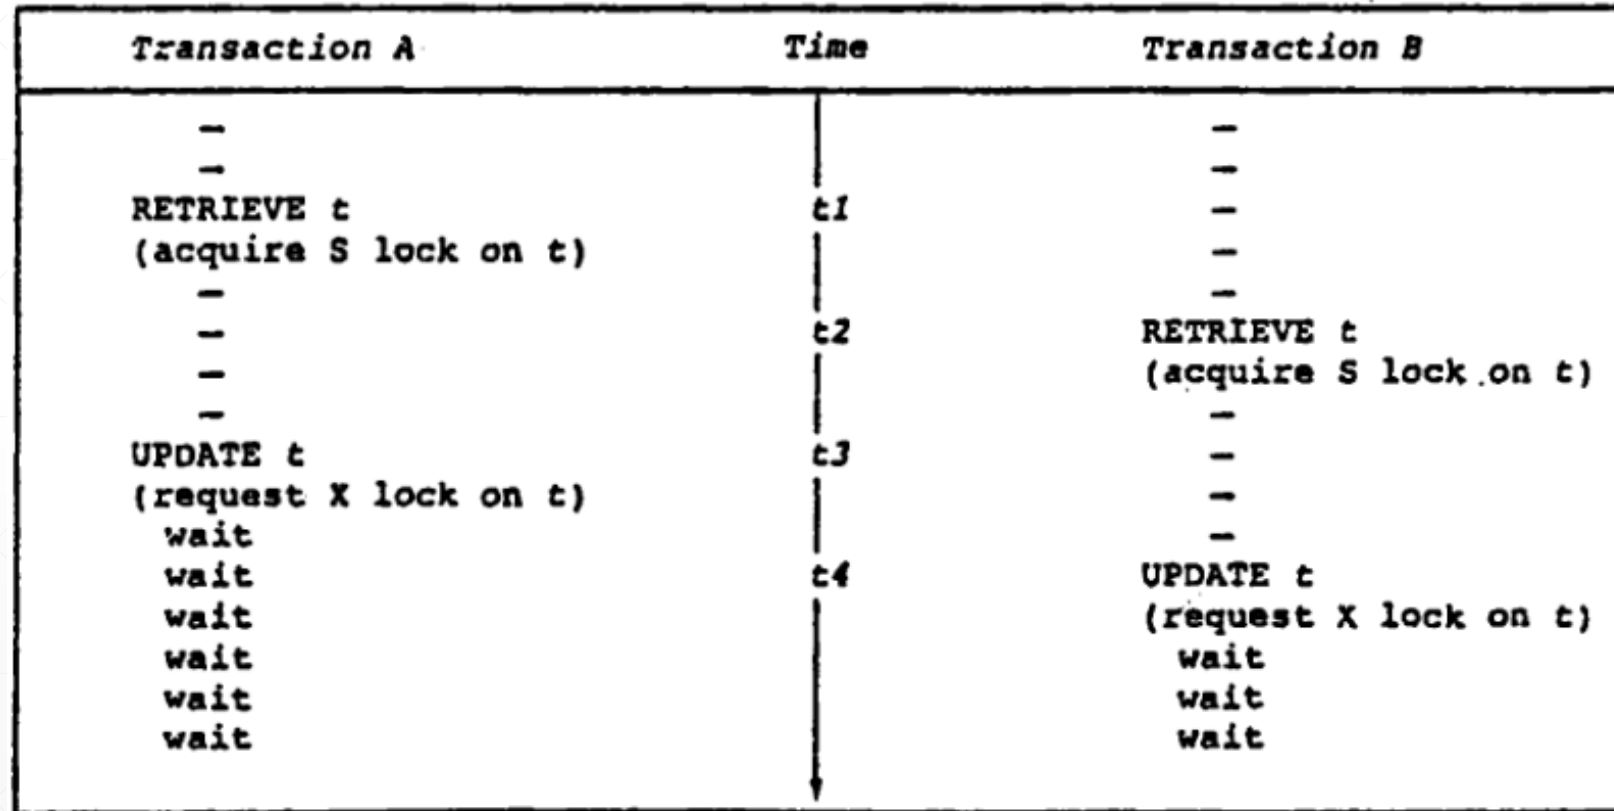


Fig. 16.6 No update is lost, but deadlock occurs at time t4

The Uncommitted Dependency Problem




Fig. 16.2 Transaction A becomes dependent on an uncommitted change at time t2

An uncommitted dependency occurs when a second transaction relies on a change which has not yet been committed, which is rolled back after the second transaction has begun.

The Uncommitted Dependency Problem - Revisited

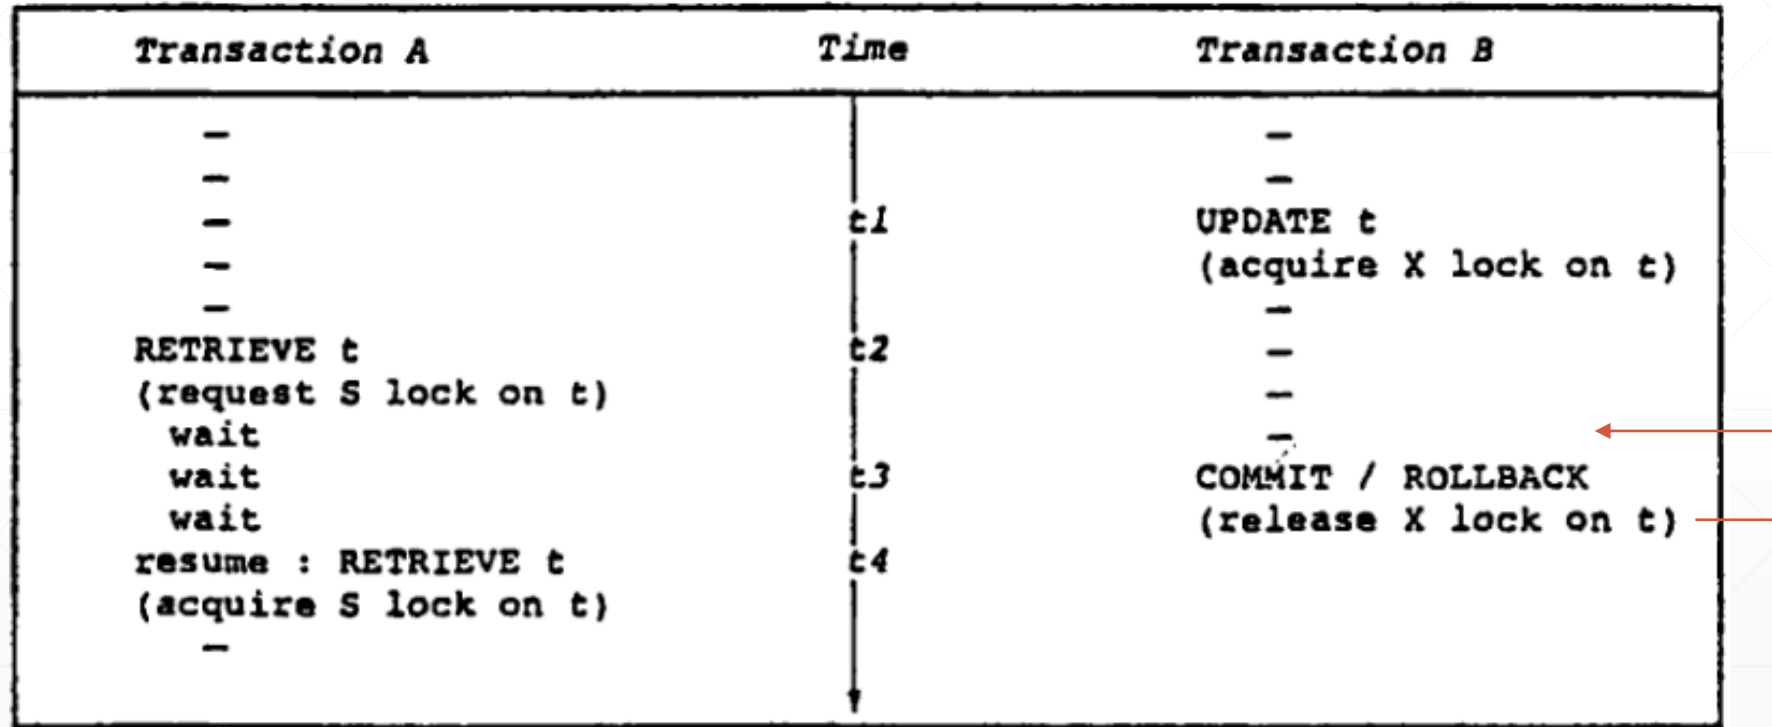
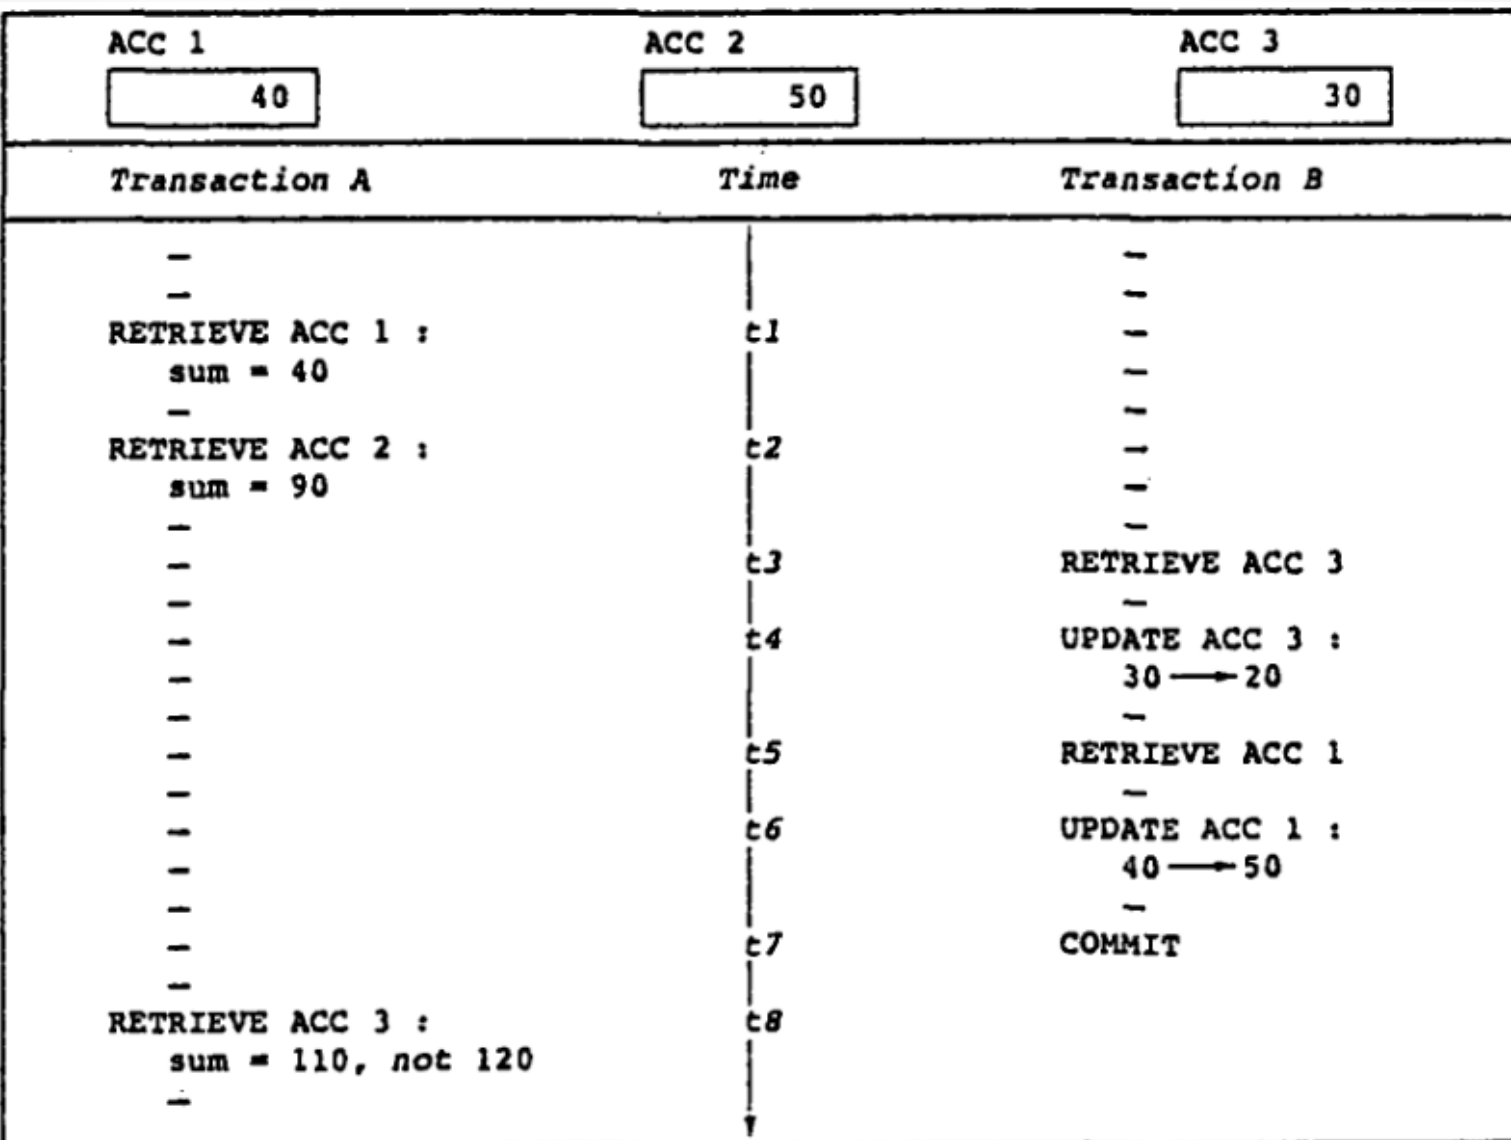


Fig. 16.7 Transaction A is prevented from seeing an uncommitted change at time t2



The Inconsistent Analysis Problem

An inconsistent analysis occurs when totals are calculated during interleaved updates.

Fig. 16.4 Transaction A performs an inconsistent analysis

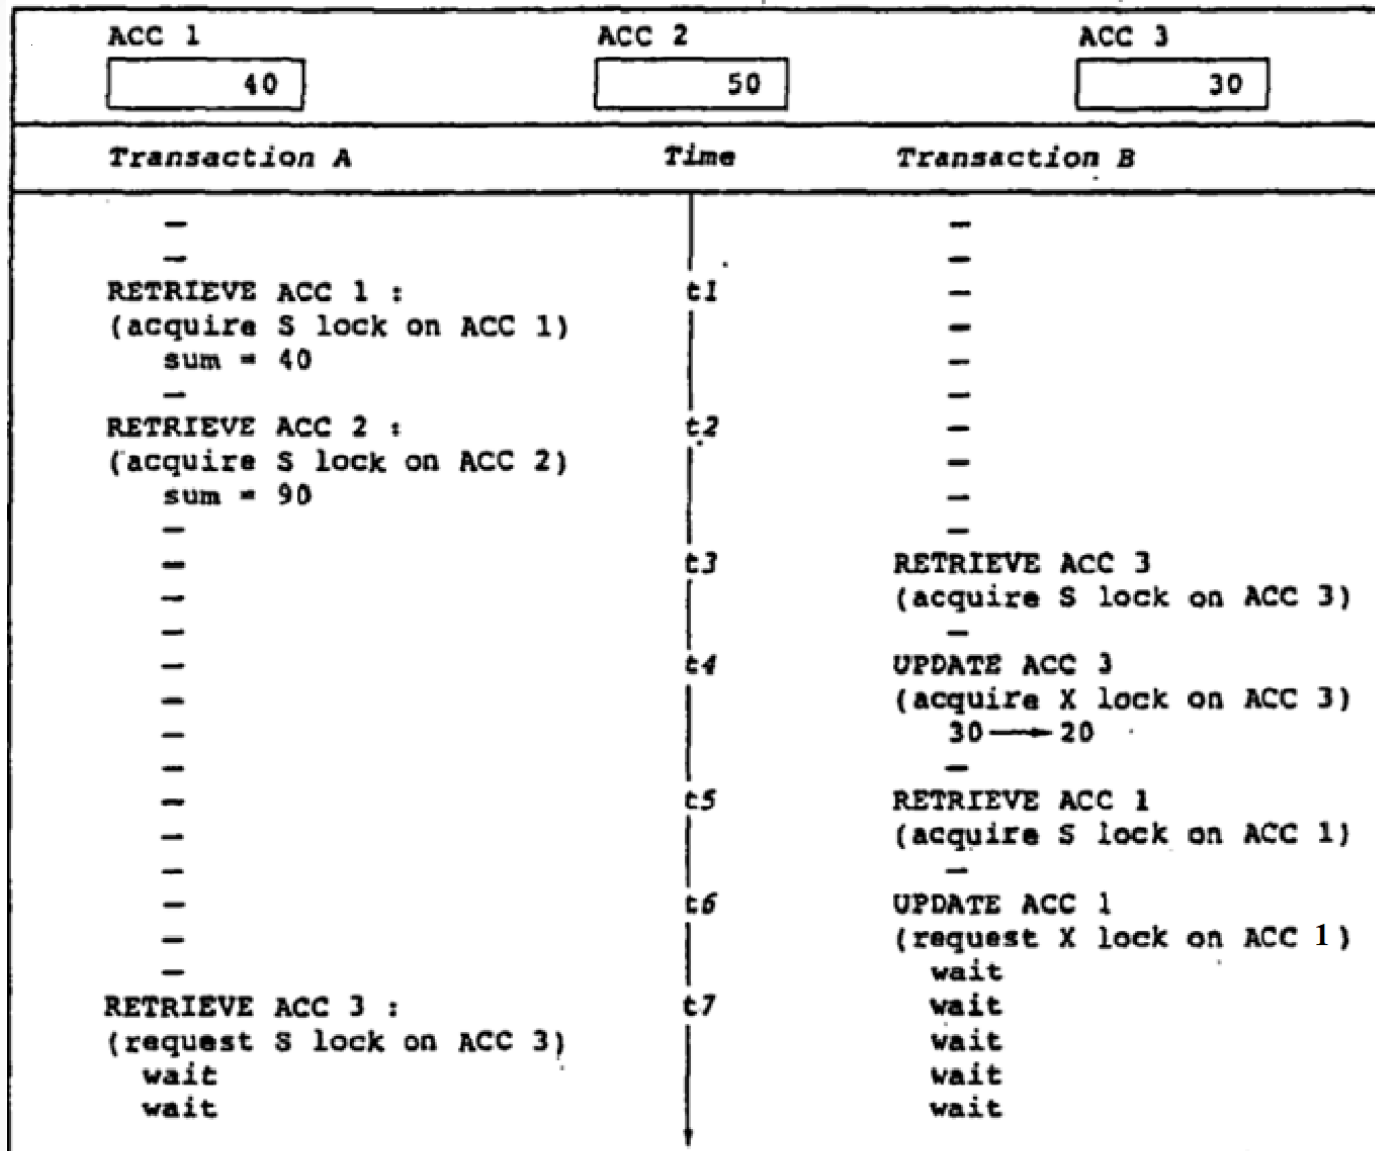


Fig. 16.9 Inconsistent analysis is prevented, but deadlock occurs at time t7

The Inconsistent Analysis Problem - Revisited

Extensions to Basic Two-Phase Locking Protocol

- **Two-phase locking:** Allows a transaction to lock a new data item only if that transaction has not yet unlocked any data item.
 - Not free from cascading rollbacks
- **Strict two-phase locking:** A transaction must hold all its exclusive locks till it commits/aborts.
 - Ensures recoverability of freedom from cascading rollback.
- **Rigorous two-phase locking:** A transaction must hold all locks till commit/abort.
 - Transactions can be serialized in the order in which they commit.
 - Ensures recoverability of freedom from cascading rollback.

Deadlock Handling

- Various locking protocols do not guard against deadlocks.
- System is **deadlocked** if there is a set of transactions such that every transaction in the set is waiting for another transaction in the set.

T_3	T_4
lock-X(B) read(B) $B := B - 50$ write(B)	lock-S(A) read(A) lock-S(B)
lock-X(A)	

Deadlock Prevention Strategies

- The locking protocol may be modified to avoid deadlock by using Wait-Die Scheme and Wound-Wait Scheme
- **Wait-Die:** Transaction 2 waits if it is older than 1; otherwise it dies (rolls back)
- **Wound-Wait:** Transaction 2 wounds 1 if it is older; that is, it rolls it back
- Net effect: oldest transaction wins
- In both schemes, a rolled back transactions is restarted with its original timestamp.
 - Ensures that older transactions have precedence over newer ones, and starvation is thus avoided.

Example: Wait-Die

T1	T2
lock-S(A)	
read(A)	
	lock-S(B)
	read(B)
lock-X(B)	
	lock-X(A)

- T1 is older so it is allowed to wait.
- T2 is younger so it is roll backed.
 - Its locks are released.
 - Allows T1 to proceed.

Example: Wound-Wait

T1	T2
lock-S(A)	
read(A)	
	lock-S(B)
	read(B)
lock-X(B)	

- T1 is older so T2 is roll backed and that allows to T1 proceed.

Deadlock Prevention Strategies (Cont.)

Wait-Die Scheme — non-preemptive

- Older transaction may wait for younger one to release data item.
- Younger transactions never wait for older ones; they are rolled back instead.
- A transaction may die several times before acquiring a lock.

Wound-Wait Scheme — preemptive

- Older transaction *wounds* (forces rollback) of younger transaction instead of waiting for it.
- Younger transactions may wait for older ones.
- Fewer rollbacks than *wait-die* scheme.

Deadlock Prevention Strategies (Cont.)


- **Timeout-Based Schemes:**

- A transaction waits for a lock only for a specified amount of time. After that, the wait times out and the transaction is rolled back.
- Ensures that deadlocks get resolved by timeout if they occur.
- Simple to implement.
- But may roll back transaction unnecessarily in absence of deadlock
 - Difficult to determine good value of the timeout interval.
- Starvation is also possible.


Deadlock Detection

- If deadlocks are not prevented, the system must deal with them by using a deadlock detection and recovery scheme.
- **Wait-for graph**
 - *Vertices*: transactions
 - *Edge from $T_i \rightarrow T_j$* : if T_i is waiting for a lock held in conflicting mode by T_j
- The system is in a deadlock state if and only if the wait-for graph has a cycle.
- Invoke a deadlock-detection algorithm periodically to look for cycles.

Deadlock Detection (Cont.)



Wait-for graph without a cycle



Wait-for graph with a cycle

Deadlock!

Isolation Levels

- The isolation level indicates the degree of interaction that is allowed from other transactions during the execution of transaction.
 - **Read uncommitted:** Even uncommitted records may be read.
 - **Read committed:** Only committed records can be read, but successive reads of record may return different (but committed) values.
 - **Repeatable read:** Ensures that a read is repeatable throughout the transaction.
 - **Serializable:** Usually ensures serializable execution.
- The higher the level of isolation, the less interference, and the lower concurrency.

Dirty Read, Nonrepeatable Read and Phantom

- **Dirty read:** A transaction reads values written by another transaction that hasn't committed yet.
- **Nonrepeatable read:** A transaction reads the same object twice during execution and finds a different value the second time, although the transaction has not changed the value in the meantime.
- **Phantom read:** A transaction re-executes a query returning a set of rows that satisfy a search condition and finds that the set of rows satisfying the condition has changed as a result of another recently committed transaction.

Isolation Levels and Problems

ISOLATION LEVEL	DIRTY READ	NONREPEATABLE READ	PHANTOM READ
READ UNCOMMITTED	Y	Y	Y
READ COMMITTED	N	Y	Y
REPEATABLE READ	N	N	Y
SERIALIZABLE	N	N	N

References

- A. Silberschatz, HF. Korth, S. Sudarshan, Database System Concepts, 7th Ed., McGraw-Hill, 2019.
 - Chapter 18, <https://www.db-book.com/db7/slides-dir/PPTX-dir/ch18.pptx> (modified)
- Edward Sciore, Database Design and Implementation, Wiley, 2009.
 - Chapter 14
- C.J. Date, An Introduction to Database Systems, 8th Ed., 2003.
 - Chapter 16